

BLG453E COMPUTER VISION
Fall 2017 Term

İTÜ 
Istanbul Technical University
Computer Engineering Department

Instructor: Prof. Gözde ÜNAL

Teaching Assistant: Enes ALBAY

Week 1-2 Topics and Learning Objectives:

Introduction to Computer Vision

Pointwise Image Processing

Image Intensity Transformations, Image Histograms,

Image Enhancement

At the end of Week 1-2: Students will be able to:

1. Discuss the main problems of computer (artificial) vision, its uses and applications
2. Design and implement various image transforms: point-wise transforms, neighborhood operation-based spatial filters, and geometric transforms over images

What is an image?

An image is multi-dimensional signal that measures a physical quantity.

Multi-dimensional

- 2D: Image (as a function of x and y)
- 3D: Video (as a function of x , y , and t)
- (others, e.g. CT, MRI)

Physical quantity

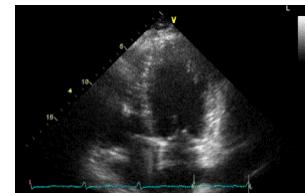
- Typically brightness for standard photographs
- Many other possibilities (temperature, depth, acoustic properties, etc.)



Digital photo



Thermal image



Ultrasound video

Multi-dimensional function

Mathematically, an image is a multi-dimensional function. For example, a grayscale image can be expressed as a mapping from the set of real numbers in 2D space to the set of real numbers (brightness).

We could write

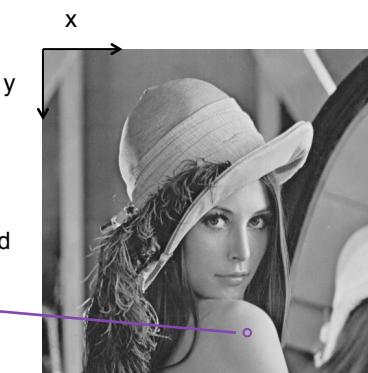
$$I : \mathbb{R}^2 \rightarrow \mathbb{R}$$

Or simply

$$I(x, y)$$

Where x and y are the spatial variables, and I is the intensity

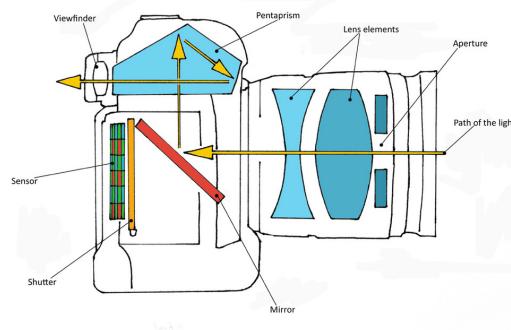
$$I(200, 400) = 178$$



Camera

Typical components

- Aperture to let light in
- Lens
- Photosensitive image plane
- Housing to keep stray light out



<http://murrayparkphotography.weebly.com/inside-a-dslr.html>

Pinhole camera

- All light passes through a single point, known as the *centre of projection*.
- On the other side of the camera is film that captures striking light.
- This creates a *perspective projection*.

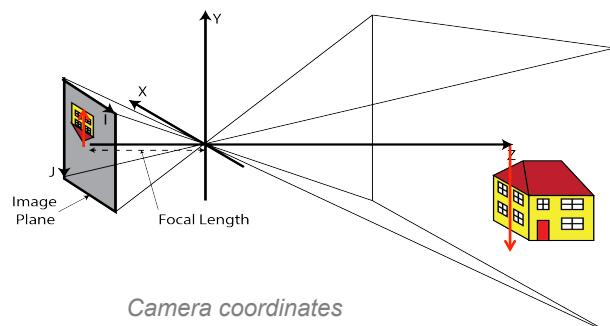


<http://petapixel.com/2013/03/26/how-to-create-a-homemade-large-format-pinhole-camera-using-a-shoebox/>

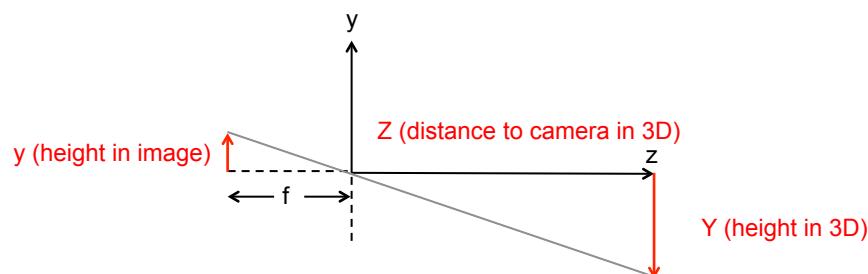
Pinhole camera model

This can be modelled mathematically

- Whilst simple, the model is reasonably realistic
- Deviations occur due to lenses that create distortions



Pinhole camera model

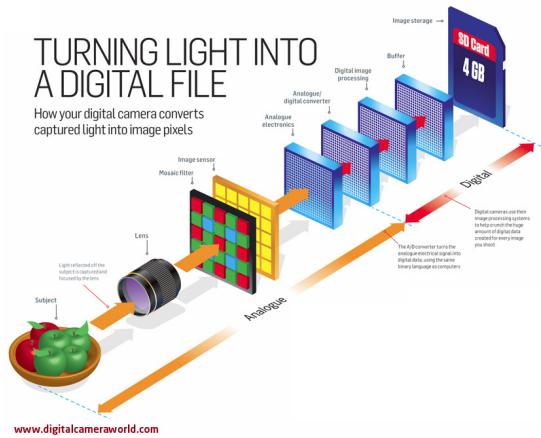


$$\text{By similar triangles, } \frac{y}{f} = \frac{Y}{Z} \quad \text{or} \quad y = \frac{fY}{Z}$$

⇒ As Z increases, y decreases (i.e., farther away objects appear smaller!)

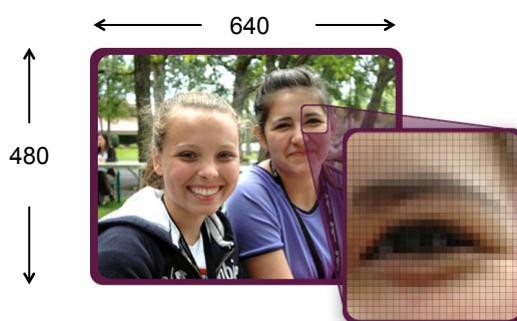
From light to pixels

The light impinging on the image sensor is a continuous signal. This is converted to a *digital* signal through *sampling* and *quantisation*, to form a set of *pixels* comprising the image.



Digital image

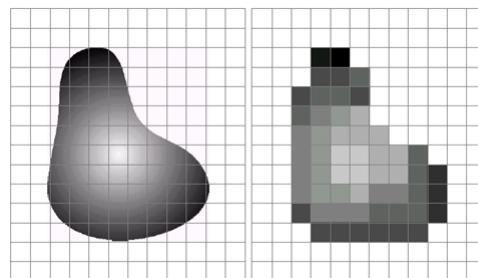
A digital image is a collection of pixels, arranged on a 2D grid. Each pixel stores colour information comprising the image. The image *resolution* is the size of the image, in pixels (in width and height).



<http://www.g2cs.org/media/interactives/image-compression/00.html>

Sampling

Sampling converts a continuous signal into a discrete one. The light impinging the image plane is a continuous signal in x and y. In a digital camera, this signal is sampled on a regular grid. The sampling rate determines the resolution.



Well-known devices

iPhone 6



Rear camera
3264 x 2448 ~ 8 MP

Samsung Galaxy S6



Rear camera
5312 x 1988 ~ 16 MP

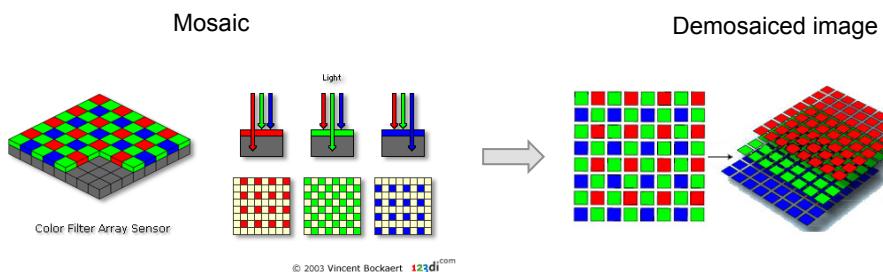
HDTV



display
1920x1080 ~ 2 MP

Colour filter array and demosaicing

Light captured by digital cameras is usually binned into separate red, green, and blue values using a *colour filter array*, often made of a GRBG pattern called the Bayer pattern. A *demoslicing* algorithm interpolates the data so that each pixel has a red, green, and blue (RGB) value.



<https://www.youtube.com/watch?v=2-stCNB8jT8>

Quantisation

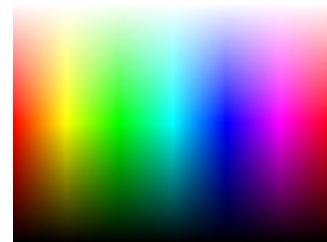
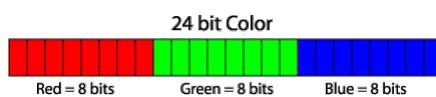
Quantisation limits the values a pixel can have to a finite set. For example, in a greyscale image, one normally uses 8 bits per pixel, so there are $2^8 = 256$ different intensities, normally represented in the range [0, 255].



<https://songbirdsings13.wordpress.com/year-12/new-media/unit-16/pmd1-unit-19/>

Quantisation

In standard colour photographic images, one uses 8 bits for *each* colour channel (red, green, blue), or 24 bits per pixel. That means there are 2^{24} possible colours for a pixel. This is roughly 16.7 million colours!



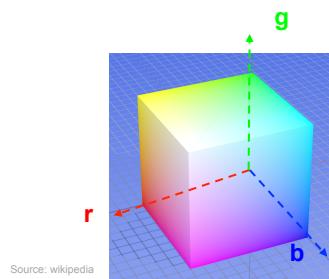
Colour

Colour images typically have three colour *channels* corresponding to the amount of red, green, and blue present at each pixel. Combining them together produces the final colour.



RGB colour space

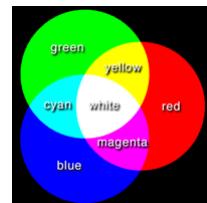
- Has its origin in colour television, now used in displays (flat panels, phones)
- *Additive* mixing or red, green, and blue light form the final colour
- RGB is a colour space
 - Red axis, green axis, blue axis
 - Values typically in the range from 0 (none) to 255 (full colour) along each axis
 - A colour is a point in this space, represented as a vector $[r, g, b]^T$



RGB colour

RGB uses additive colour mixing, because it describes what kind of *light* needs to be *emitted* to produce a given colour. Light is added together to create form from out of the darkness.

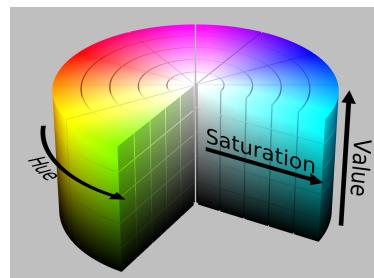
Colour Component			Colour Common Name
R	G	B	
0	0	0	Black
0	0	255	Blue
0	255	0	Green
0	255	255	Cyan
255	0	0	Red
255	0	255	Magenta
255	255	0	Yellow
255	255	255	White



⇒ Of course, one can have values, like [255, 127, 0] -- orange

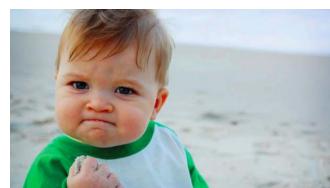
HSV

- Transforms the colour space to be more intuitive and perceptually relevant than RGB, using a cylindrical coordinate system.
 - H (hue) corresponds to the discernible colour based on the dominant wavelength
 - S (saturation) corresponds to the vividness of the colour. Low saturation means a grayscale colour in the centre of the cylinder.
 - V (value) corresponds to brightness.
- ⇒ For example, a bright red colour has a red hue, full saturation, and high value.



Other color spaces exist such as CIE LAB. In this course, we'll work in RGB color space

HSV



Original image



H



S



V

Other common image types

RGBD

- In addition to a colour image (RGB), acquires a depth image (D), which represents the distance from each pixel from the camera







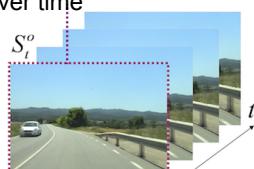
Volumetric images

- Image data stored as voxels (small cubes, or 3D pixels). Examples include computed tomography (CT) or magnetic resonance (MRI)



Video

- A sequence of images over time



Point Processing of Images



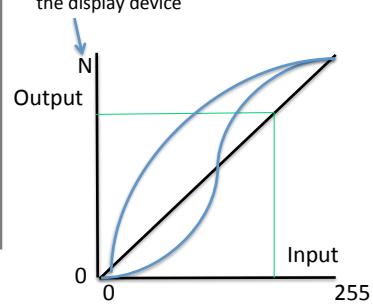





- gamma
- brightness
original
+ brightness
+ gamma

- In a digital image, point = pixel.
- Point processing transforms a pixel's intensity value as function of its value alone;
- it does not depend on the values of the pixel's neighbors.

Image: a 2D pattern of intensity values

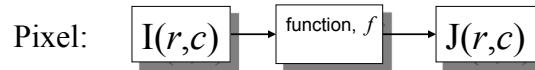
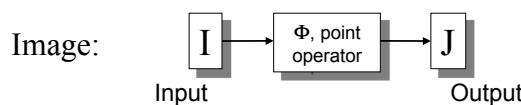
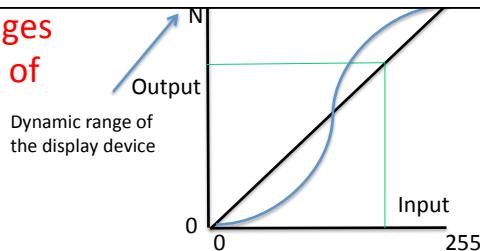


Point Processing of Images

- Brightness and contrast adjustment
- Gamma correction
- Histogram equalization
- Histogram matching

19/09/17

Point Operations on Images via Functional Mappings of intensities



The transformation of image I into image J is accomplished by replacing each input intensity, g , with a specific output intensity, k , at every location (r,c) where $I(r,c) = g$.

If $I(r,c) = g$
and $f(g) = k$
then $J(r,c) = k$.

The rule that associates k with g is usually specified with a function, f , so that $f(g) = k$.

Point Operations on Images via Functional Mappings

One-band Image

$$J(r,c) = f(I(r,c)), \text{ for all pixels locations } (r,c).$$

Three-band Image

$$J(r,c,b) = f(I(r,c,b)), \text{ or } J(r,c,b) = f_b(I(r,c,b)), \text{ for } b = 1,2,3 \text{ and all } (r,c).$$

1999-2007 by Richard Alan Peters II

Point Operations on Images via Functional Mappings

One-band Image

Either all 3 bands
are mapped through
the same function,
 f , or ...

$$J(r,c) = f(I(r,c)), \text{ for all pixels locations } (r,c).$$

Three-band Image

$$J(r,c,b) = f(I(r,c,b)), \text{ or } J(r,c,b) = f_b(I(r,c,b)), \text{ for } b = 1,2,3 \text{ and all } (r,c)$$

... each band is
mapped through
a separate func-
tion, f_b .

Point Operations Using Look-up Tables

A look-up table (LUT)
implements a
functional mapping.

If $k = f(g)$,
for $g = 0, \dots, 255$,
and if k takes on
values in $\{0, \dots, 255\}$, ...

... then the LUT
that implements f
is a 256×1 array
whose $(g+1)^{\text{th}}$
value is $k = f(g)$.

To remap a 1-band
image, I , to J :

$$J = \text{LUT}(I)$$

Point Operations Using Look-up Tables

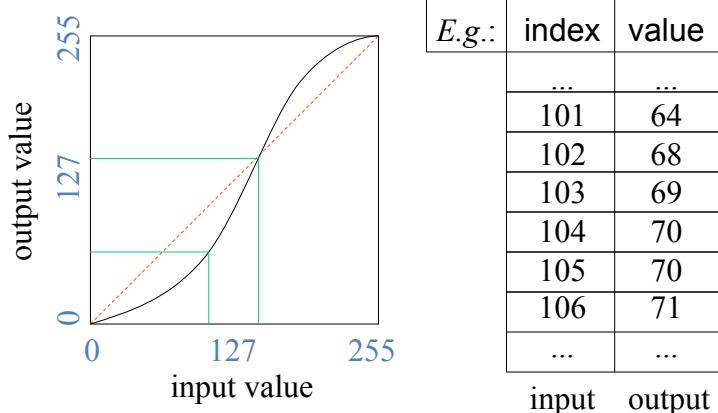
If I is 3-band, then

- a) each band is mapped separately using the same LUT for each band *or*
- b) each band is mapped using different LUTs – one for each band.

a) $J = \text{LUT}(I)$, or

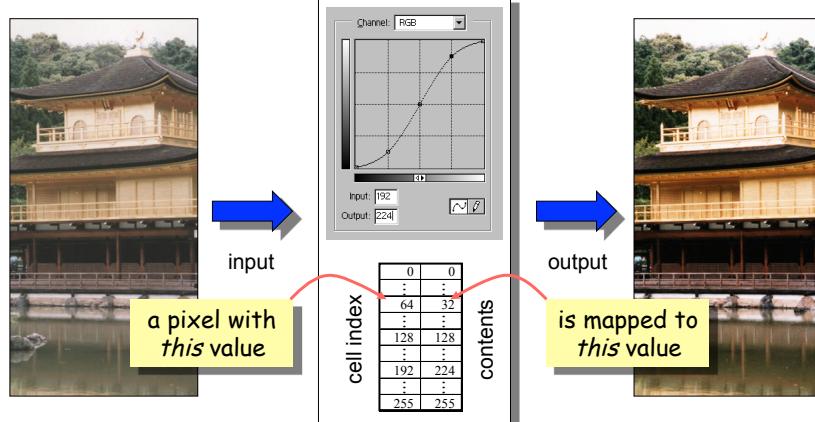
b) $J(:, :, b) = \text{LUT}_b(I(:, :, b))$ for $b = 1, 2, 3$.

Point Operations = Look-up Table Ops



19/09/17

Look-Up Tables



19/09/17

33

by Richard Alan Peters II

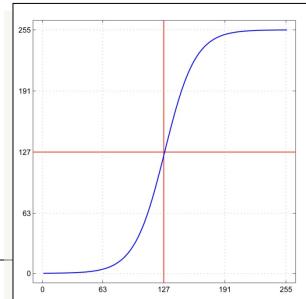
How to Generate a Look-Up Table

For example:

Let $a = 2$.

Let $x \in \{0, \dots, 255\}$

$$\sigma(x; a) = \frac{255}{1 + e^{-a(x-127)/32}}$$



Or in Matlab:

```
a = 2;
x = 0:255;
LUT = 255 ./ (1+exp(-a*(x-127)/32));
```

This is just
one example.

19/09/17

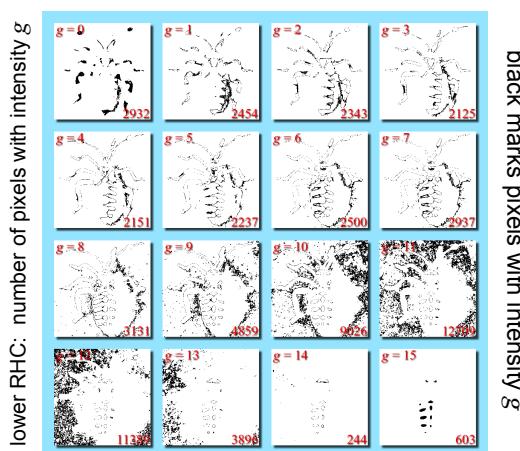
34

by Richard Alan Peters II

The Histogram of a Grayscale Image



16-level (4-bit) image



19/09/17

35

by Richard Alan Peters II

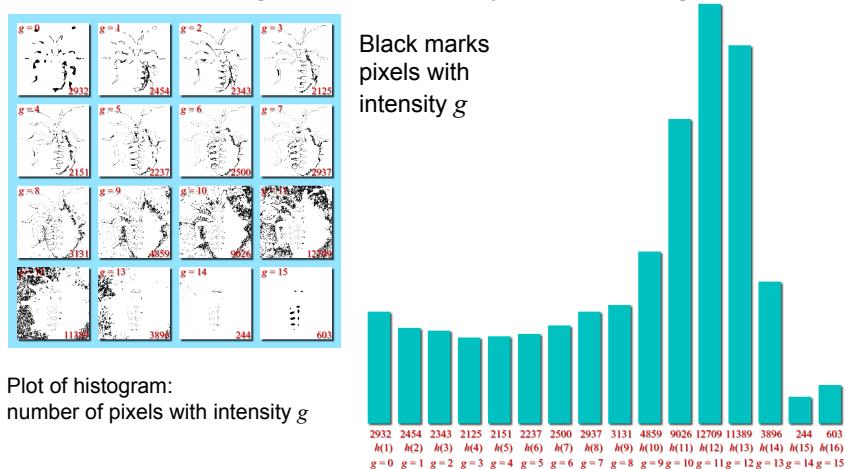
The Histogram of a Grayscale Image

- Let I be a 1-band (grayscale) image.
- $I(r,c)$ is an 8-bit integer between 0 and 255.
- Histogram, h_I , of I :
 - a 256-element array, h_I
 - $h_I(g)$, for $g = 0, 2, 3, \dots, 255$, is an integer
 - $h_I(g) = \text{number of pixels in } I \text{ that have value } g$.

19/09/17

36

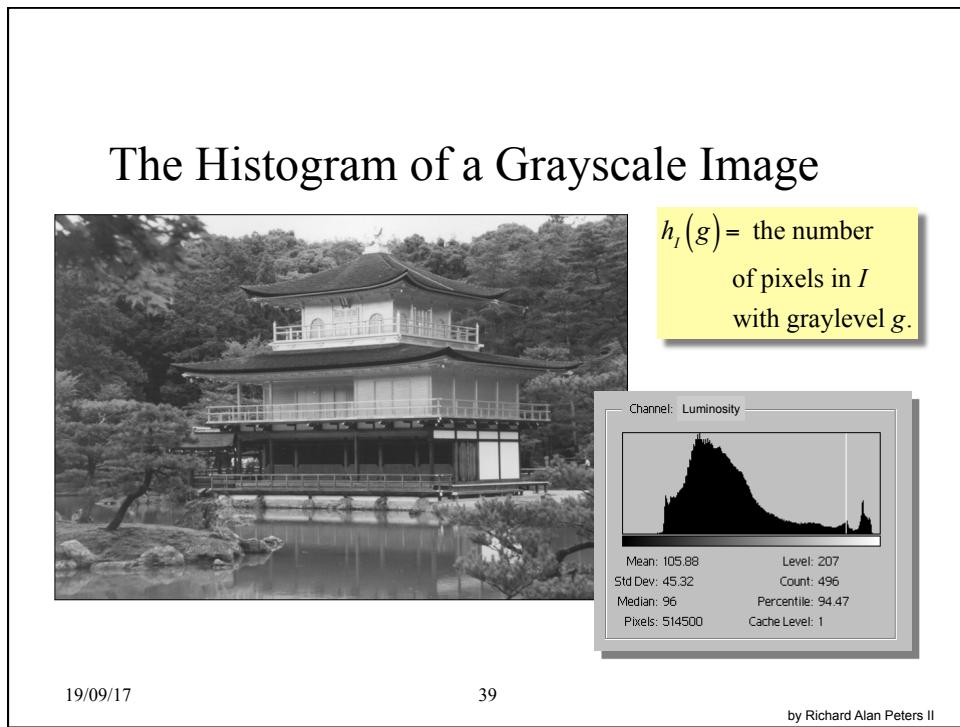
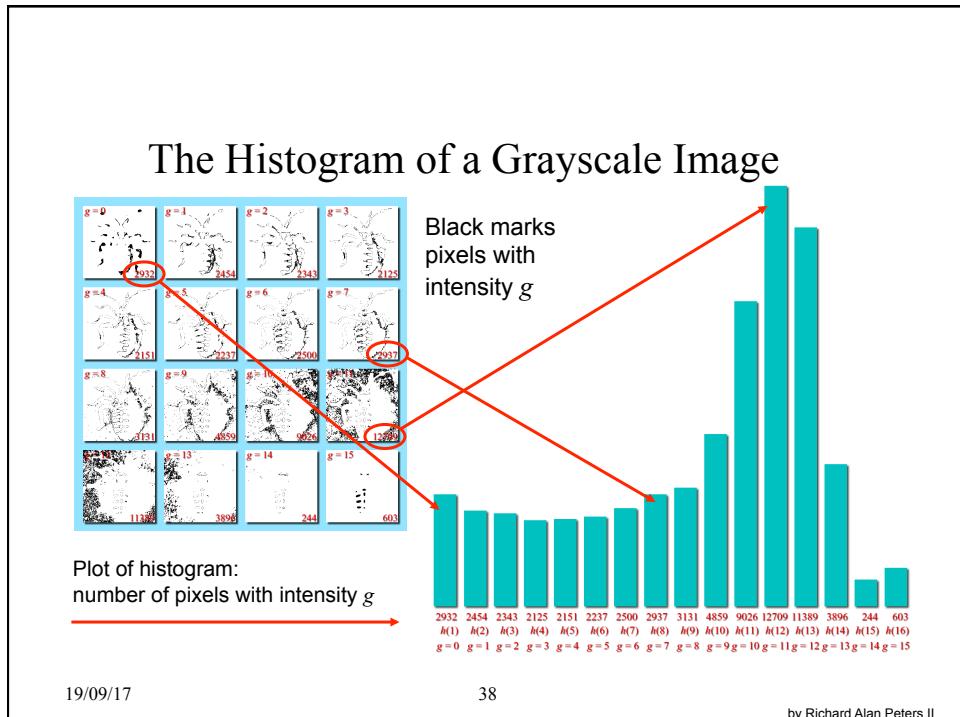
The Histogram of a Grayscale Image



19/09/17

37

by Richard Alan Peters II



The Histogram of a Color Image

- If I is a 3-band image (truecolor, 24-bit)
- then $I(r,c,b)$ is an integer between 0 and 255.
- Either I has 3 histograms:
 - $h_R(g) = \#$ of pixels in $I(:,:,1)$ with intensity value g
 - $h_G(g) = \#$ of pixels in $I(:,:,2)$ with intensity value g
 - $h_B(g) = \#$ of pixels in $I(:,:,3)$ with intensity value g
- or 1 vector-valued histogram, $h(g, 1, b)$ where
 - $h(g, 1, 1) = \#$ of pixels in I with red intensity value g
 - $h(g, 1, 2) = \#$ of pixels in I with green intensity value g
 - $h(g, 1, 3) = \#$ of pixels in I with blue intensity value g

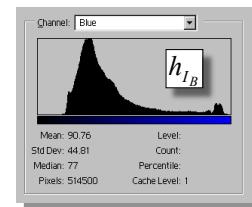
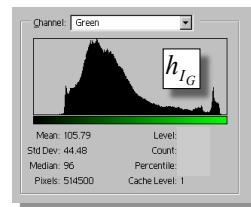
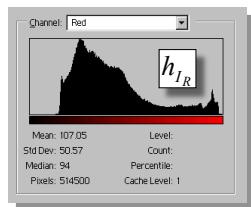
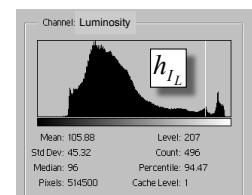
19/09/17

40

by Richard Alan Peters II

The Histogram of a Color Image

There is one histogram per color band R, G, & B. Luminosity histogram is from 1 band = $(R+G+B)/3$

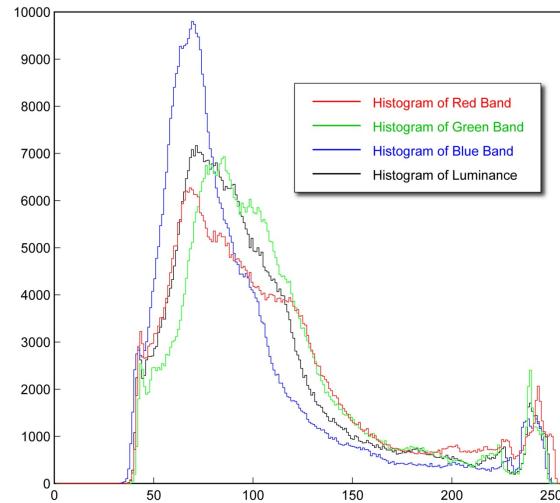


19/09/17

41

by Richard Alan Peters II

The Histogram of a Color Image



19/09/17

42

by Richard Alan Peters II

Value or Luminance Histograms

The value histogram of a 3-band (truecolor) image, I , is the histogram of the value image,

$$V(r,c) = \frac{1}{3}[R(r,c) + G(r,c) + B(r,c)]$$

Where R , G , and B are the red, green, and blue bands of I .
The luminance histogram of I is the histogram of the luminance image,

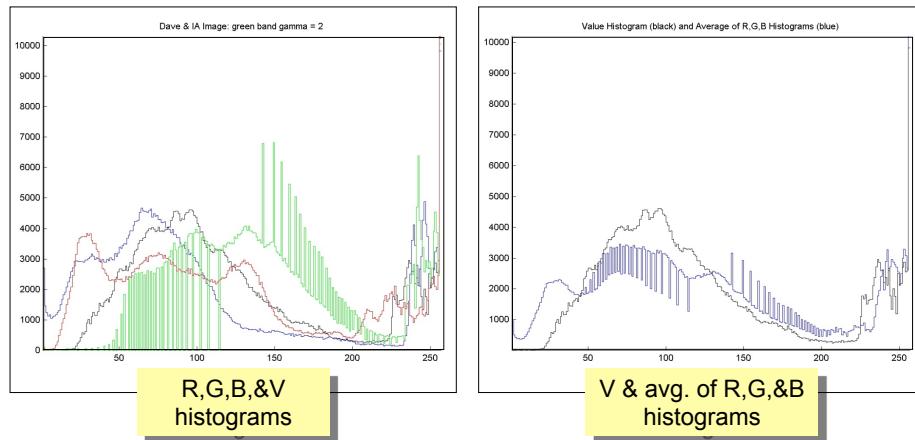
$$L(r,c) = 0.299 \cdot R(r,c) + 0.587 \cdot G(r,c) + 0.114 \cdot B(r,c)$$

19/09/17

43

by Richard Alan Peters II

Value Histogram vs. Average of R,G,&B Histograms



19/09/17

44

by Richard Alan Peters II

Multi-Band Histogram Calculator in Matlab

```
% Multi-band histogram calculator
function h=histogram(I)

[R C B]=size(I);

% allocate the histogram
h=zeros(256,1,B);

% range through the intensity values
for g=0:255
    h(g+1,1,:)= sum(sum(I==g)); % accumulate
end

return;
```

19/09/17

45

by Richard Alan Peters II

Multi-Band Histogram Calculator in Matlab

```
% Multi-band histogram calculator
function h=histogram(I)

[R C B]=size(I);

% allocate the histogram
h=zeros(256,1,B);

% range through the intensities
for g=0:255
    h(g+1,1,:)=sum(sum(I==g)); % accumulate
end

If B==3, then h(g+1,1,:) contains 3
numbers: the number of pixels in
bands 1, 2, & 3 that have intensity g.
```

Loop through all intensity levels (0-255)
 Tag the elements that have value g .
 The result is an $R \times C \times B$ logical array that has
 a 1 wherever $I(r,c,b) = g$ and 0's everywhere
 else.
 Compute the number of ones in each band of
 the image for intensity g .
 Store that value in the $256 \times 1 \times B$ histogram
 at $h(g+1,1,b)$.

$\text{sum}(\text{sum}(I==g))$ computes one number
 for each band in the image.

19/09/17

46

by Richard Alan Peters II

The Probability Density Function of an Image

- p_I is the **normalized histogram** of the image
 A : total number of pixels in the image.
$$p_I(g) = \frac{1}{A} h_I(g)$$
- $p_I(g)$ is the fraction of pixels in (a specific band of) an image
 that have intensity value g .
- $p_I(g)$ is the probability that a pixel randomly selected from
 the given image band has intensity value g .
- Sum of $p_I(g)$ over all g is 1 (whereas the sum of the
 histogram $h_I(g)$ over all g from 0 to 255 is equal to the
 number of pixels in the image)

19/09/17

Images and Probability

Parametric pdf's:

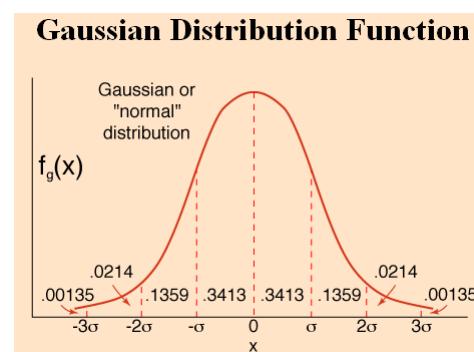
Question: Which is the probability density function that can be described by just 2 parameters, mean and variance?

Images and Probability

The distribution that can be described by just 2 parameters, mean and variance:

Important family of probability distributions: Gaussian

$$\frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$



A.k.a. Cumulative
Distribution Function.

The Probability Distribution Function of an Image

- $P_I(g)$ is the fraction of pixels in (a specific band of) an image that have intensity values less than or equal to g .
- $P_I(g)$ is the probability that a pixel randomly selected from the given band has an intensity value less than or equal to g .
- $P_I(g)$ is the cumulative (or running) sum of $p_I(g)$ from 0 through g inclusive.
- $P_I(0) = p_I(0)$ and $P_I(255) = 1$; $P_I(g)$ is nondecreasing.

Note: the Probability Distribution Function (PDF, capital letters) and the Cumulative Distribution Function (CDF) are exactly the same things. Both PDF and CDF will refer to it. However, pdf (small letters) is the *density* function.

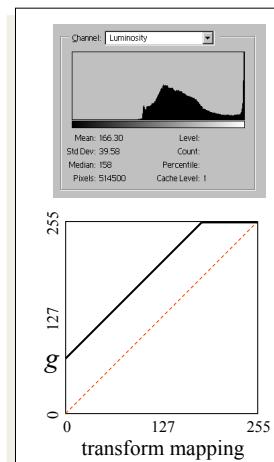
19/09/17

50

by Richard Alan Peters II



Increase Brightness



$$J_k(r,c) = \begin{cases} I_k(r,c) + g, & \text{if } I_k(r,c) + g < 256 \\ 255, & \text{if } I_k(r,c) + g > 255 \end{cases}$$

$g \geq 0$ and $k \in \{1, 2, 3\}$ is the band index.

19/09/17

51

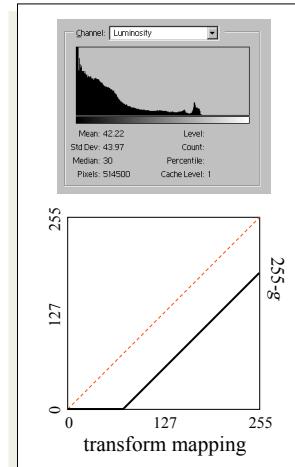
by Richard Alan Peters II

Point Processing: Decrease Brightness



$$J_k(r,c) = \begin{cases} 0, & \text{if } I_k(r,c) - g < 0 \\ I_k(r,c) - g, & \text{if } I_k(r,c) - g \geq 0 \text{ and } k \in \{1,2,3\} \end{cases}$$

is the band index.



19/09/17

52

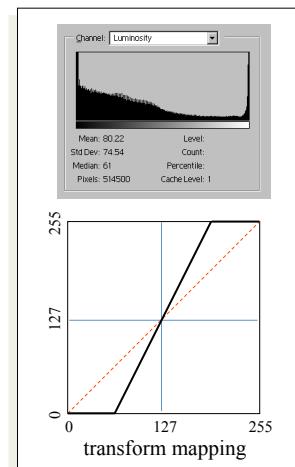
by Richard Alan Peters II

Point Processing: Increase Contrast



Let $T_k(r,c) = a[I_k(r,c) - 127] + 127$, where $a > 1.0$

$$J_k(r,c) = \begin{cases} 0, & \text{if } T_k(r,c) < 0, \\ T_k(r,c), & \text{if } 0 \leq T_k(r,c) \leq 255, \\ 255, & \text{if } T_k(r,c) > 255. \end{cases} \quad k \in \{1,2,3\}$$



19/09/17

53

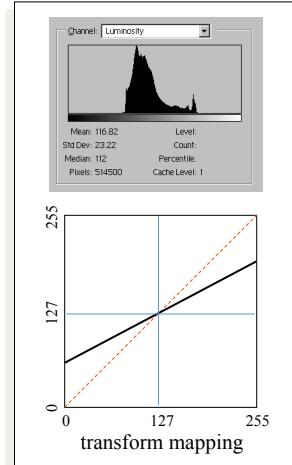
by Richard Alan Peters II

Point Processing: Decrease Contrast



$$T_k(r,c) = a[I_k(r,c) - 127] + 127,$$

where $0 \leq a < 1.0$ and $k \in \{1, 2, 3\}$.



19/09/17

54

by Richard Alan Peters II

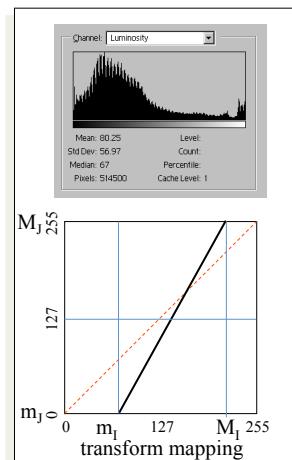
Point Processing: Contrast Stretch



Let $m_I = \min[I(r,c)]$, $M_I = \max[I(r,c)]$,
 $m_J = \min[J(r,c)]$, $M_J = \max[J(r,c)]$.

Then,

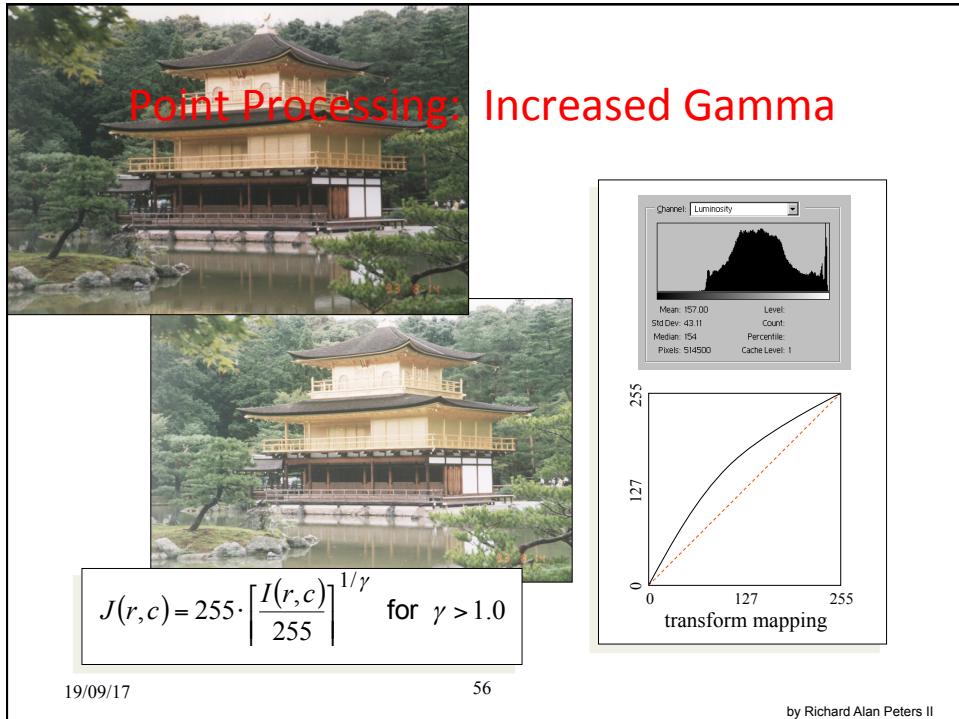
$$J(r,c) = (M_J - m_J) \frac{I(r,c) - m_I}{M_I - m_I} + m_J.$$



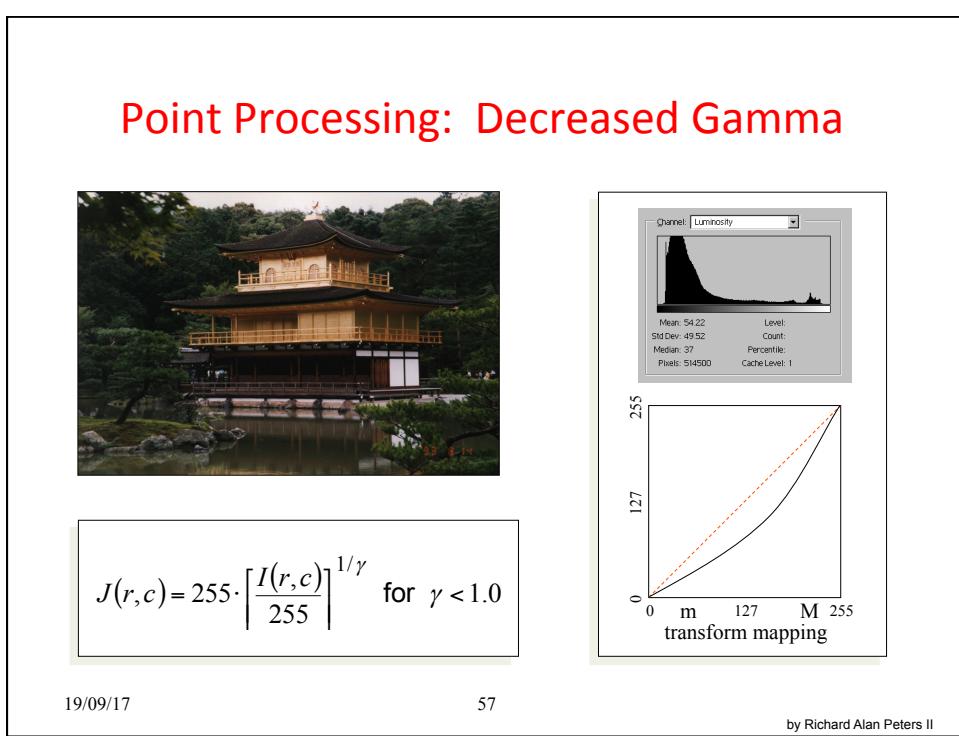
19/09/17

55

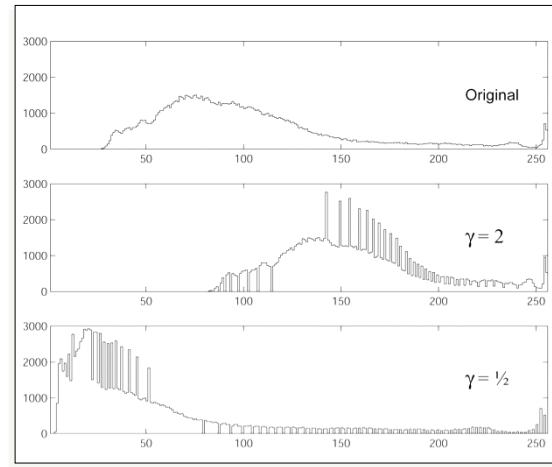
by Richard Alan Peters II



Point Processing: Decreased Gamma



Gamma Correction: Effect on Histogram

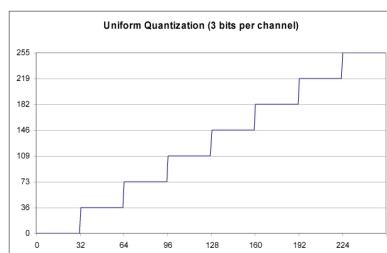
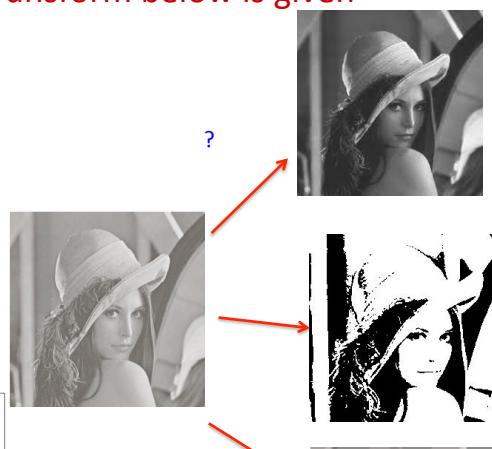
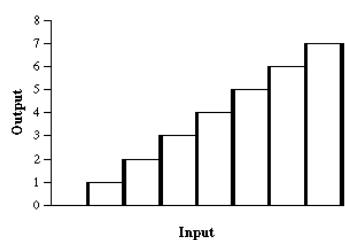


19/09/17

58

by Richard Alan Peters II

Q: Input image and the transform below is given



Uniform Quantization of Intensities is a Pointwise image transform too

Point Processing: Histogram Equalization

Recall from Probability Theory the theorem: If you plug in a random variable into its own CDF, you get a Uniform distribution: $P(x) \sim U$

Task: remap image I so that its histogram is as close to uniform as possible

Let $P_I(\gamma)$ be the cumulative (probability) distribution function of I .

Then J has, as closely as possible, the correct histogram if

$$J(r,c) = 255 \cdot P_I[I(r,c)].$$

The CDF itself is used as the LUT.

all bands
processed
similarly

19/09/17

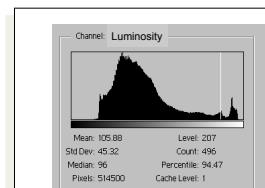
60

by Richard Alan Peters II

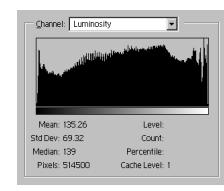
Point Processing: Histogram Equalization



$$J(r,c) = 255 \cdot P_I(g)$$



before



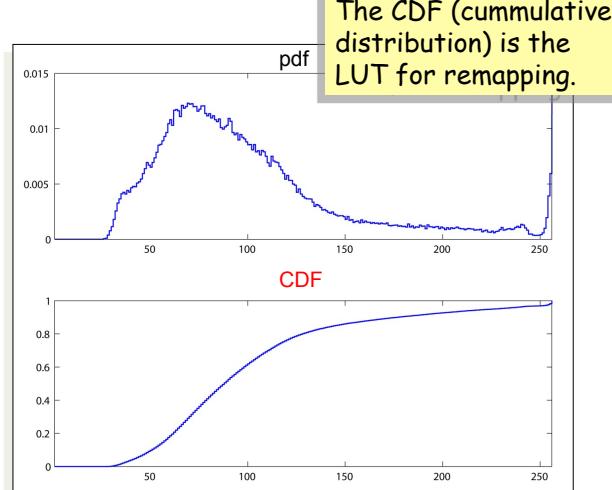
after

19/09/17

61

by Richard Alan Peters II

Histogram EQ

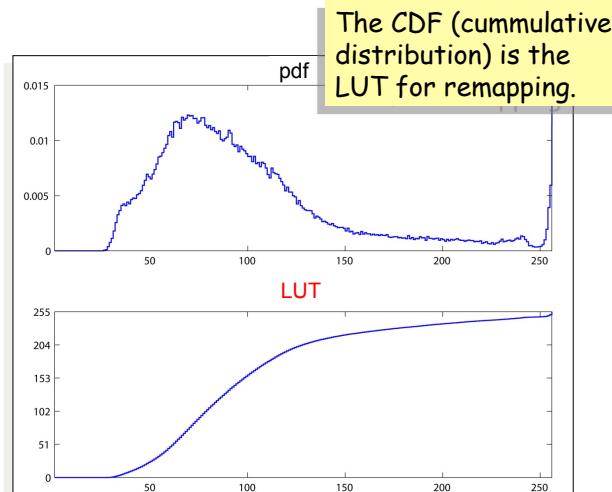


19/09/17

62

by Richard Alan Peters II

Histogram EQ

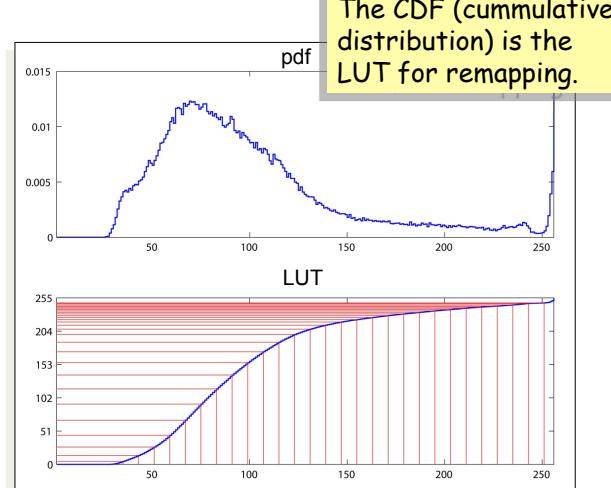


19/09/17

63

by Richard Alan Peters II

Histogram EQ

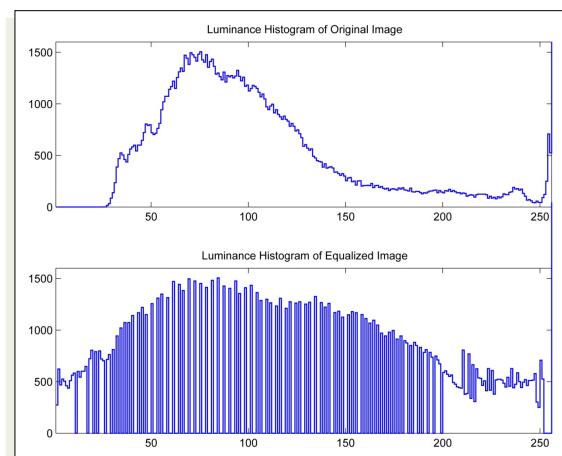


19/09/17

64

by Richard Alan Peters II

Histogram EQ



19/09/17

65

by Richard Alan Peters II

Point Processing: Histogram Matching

Task: remap image I so that it has, as closely as possible, the same histogram as image J .

Because the images are digital it is not, in general, possible to make $h_I \equiv h_J$. Therefore, $p_I \neq p_J$.

Q: How, then, can the matching be done?

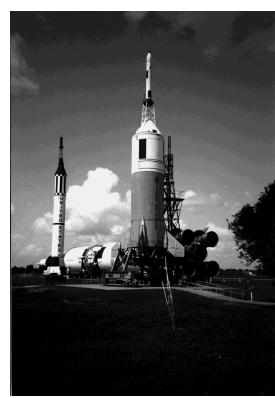
A: By matching percentiles.

19/09/17

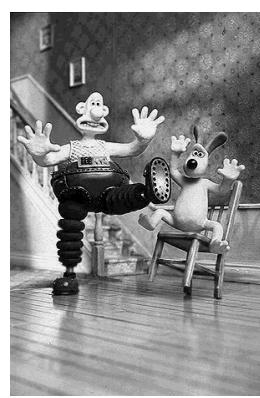
67

by Richard Alan Peters II

Example: Histogram Matching



original



target



remapped

19/09/17

68

by Richard Alan Peters II

Matching Percentiles

... assuming a 1-band image or a single band of a color image.

Recall:

- The CDF of image I is such that $0 \leq P_I(g_I) \leq 1$.
- $P_I(g_I) = c$ means that c is the fraction of pixels in I that have a value less than or equal to g_I .
- $100c$ is the *percentile* of pixels in I that are less than or equal to g_I .

To match percentiles, replace all occurrences of value g_I in image I with the value, g_J , from image J whose percentile in J most closely matches the percentile of g_I in image I .

19/09/17

69

by Richard Alan Peters II

Matching Percentiles

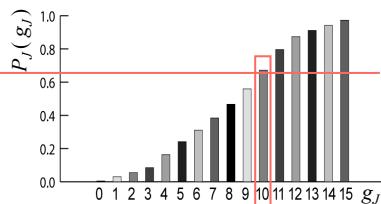
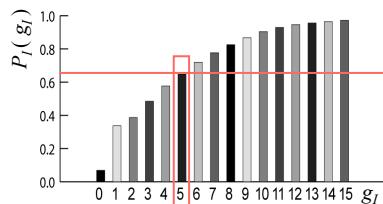
... assuming a 1-band image or a single band of a color image.

So, to create an image, K , from image I such that K has nearly the same CDF as image J do the following:

If $I(r;c) = g_I$ then let $K(r;c) = g_J$ where g_J is such that

$P_I(g_I) > P_J(g_J - I)$ AND $P_I(g_I) \leq P_J(g_J)$.

Example:
 $I(r;c) = 5$
 $P_I(5) = 0.65$
 $P_J(9) = 0.56$
 $P_J(10) = 0.67$
 $K(r;c) = 10$



19/09/17

70

by Richard Alan Peters II

Histogram Matching Algorithm

```
[R,C] = size(I) ;
K = zeros(R, C) ;
g_J = m_J;
for g_I = m_I to M_I
    while g_J < 255 AND P_I(g_I + 1) < 1 AND
          P_J(g_J + 1) < P_I(g_I + 1)
        g_J = g_J + 1;
    end
    K = K + [g_J x (I == g_I)]
end
```

... assuming a 1-band image or a single band of a color image.

This directly matches image I to image J .

$$\begin{aligned}P_I(g_I + 1) &: \text{CDF of } I \\P_J(g_J + 1) &: \text{CDF of } J \\m_J &= \min J, \\M_J &= \max J, \\m_I &= \min I, \\M_I &= \max I.\end{aligned}$$

Better to use a LUT.
See the slide (6 next)

19/09/17

71

by Richard Alan Peters II

Example: Histogram Matching

Image pdf

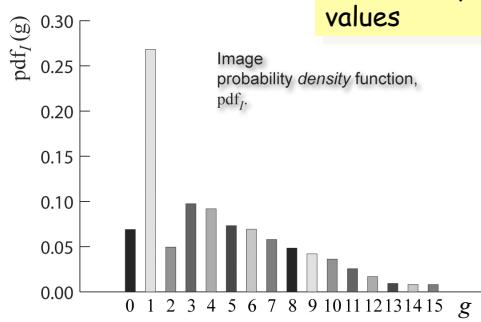
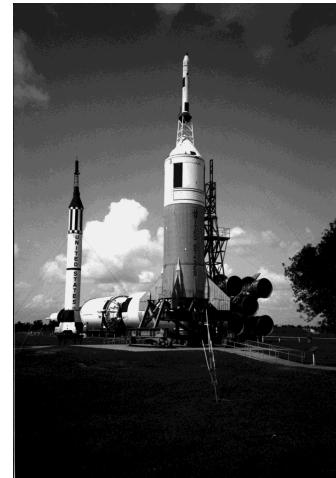


Image with 16 intensity values
Image probability density function,
 pdf_I



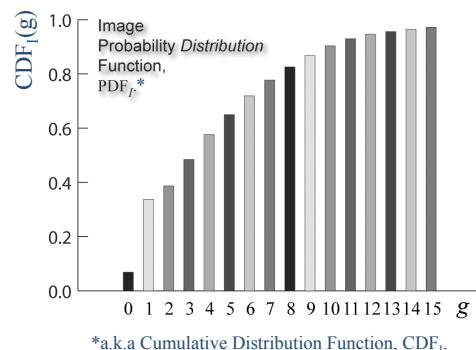
19/09/17

72

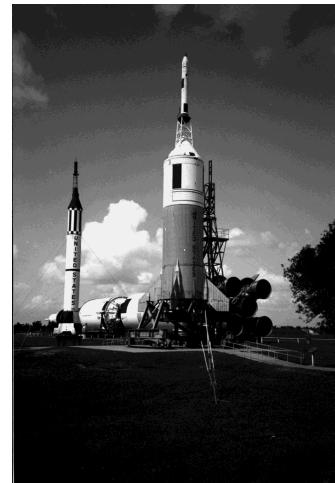
by Richard Alan Peters II

Example: Histogram Matching

Image CDF



*a.k.a Cumulative Distribution Function, CDF_I .



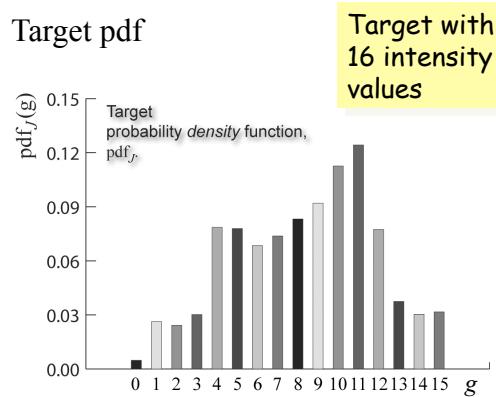
19/09/17

73

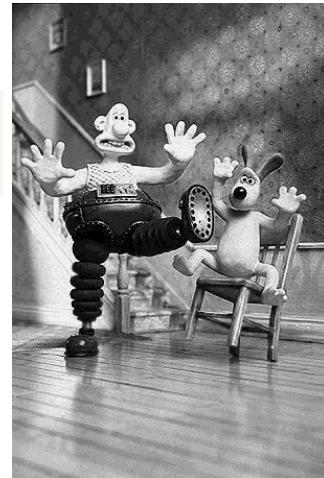
by Richard Alan Peters II

Example: Histogram Matching

Target pdf



Target with
16 intensity
values



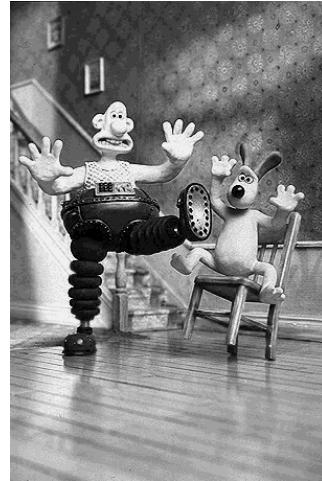
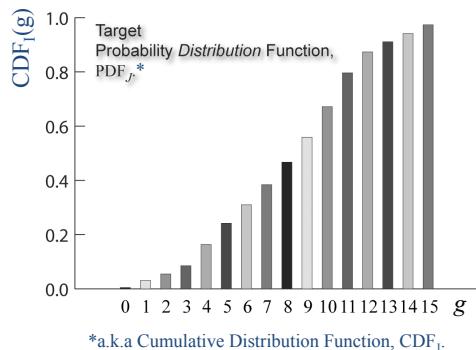
19/09/17

74

by Richard Alan Peters II

Example: Histogram Matching

Target CDF



19/09/17

75

by Richard Alan Peters II

Histogram Matching with a Lookup Table

The given algorithm (5 slides previously) matches one image to another directly. Often it is faster or more versatile to use a lookup table (LUT). Rather than remapping each pixel in the image separately, one can create a table that indicates to which target value each input value should be mapped. Then

$K = \text{LUT}[I]$ (see the graphics on the next slide)

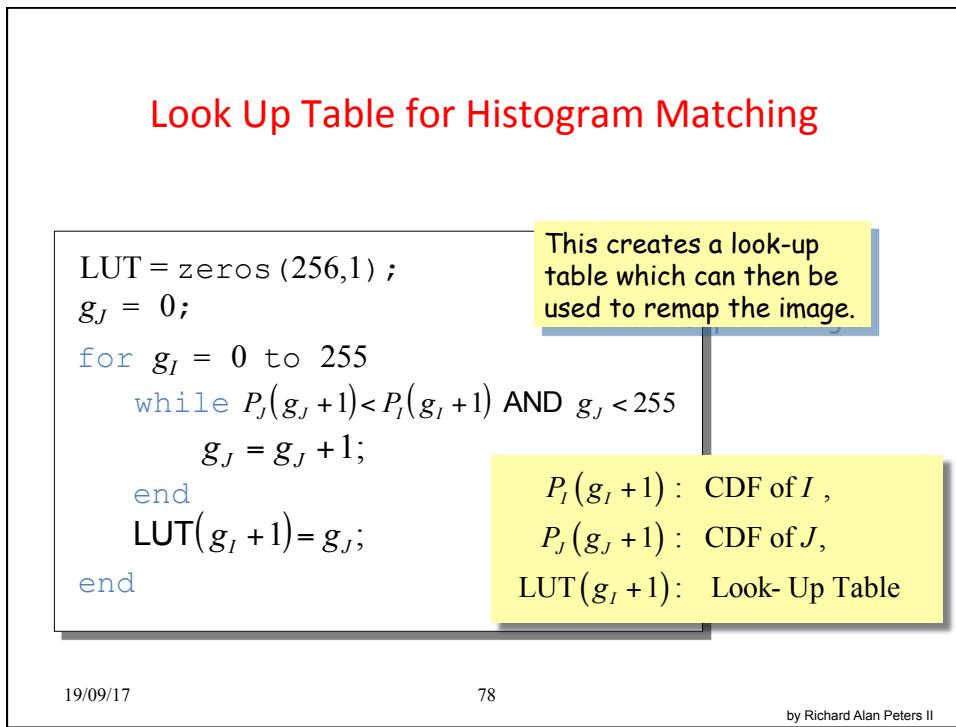
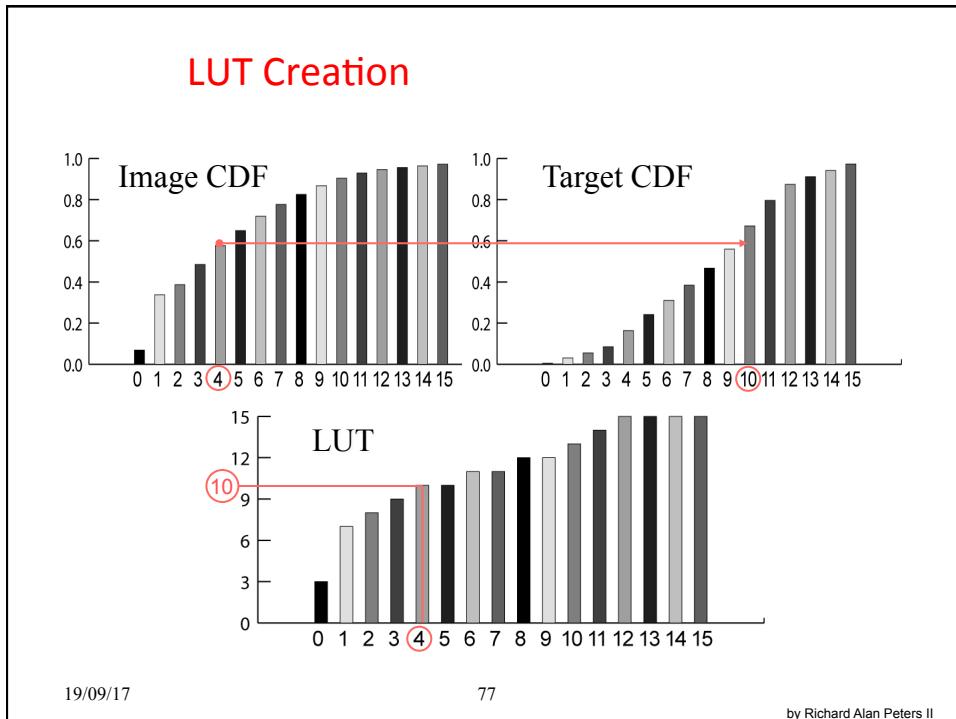
In *Matlab* if the LUT is a 256×1 matrix with values from 0 to 255 and if image I is of type **uint8**, it can be remapped with the following code:

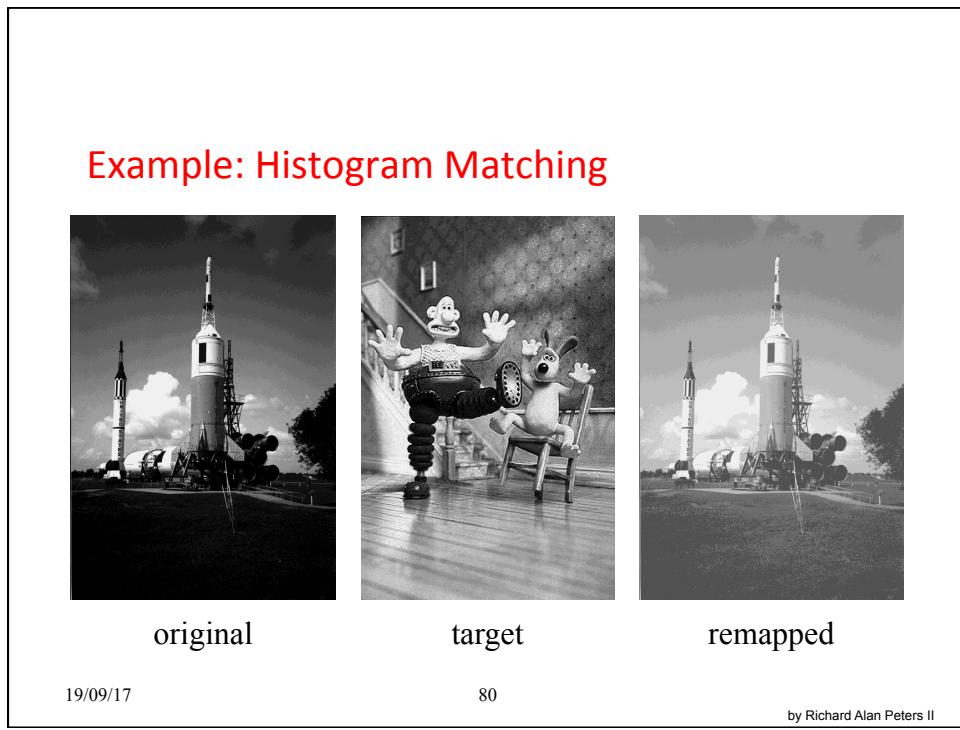
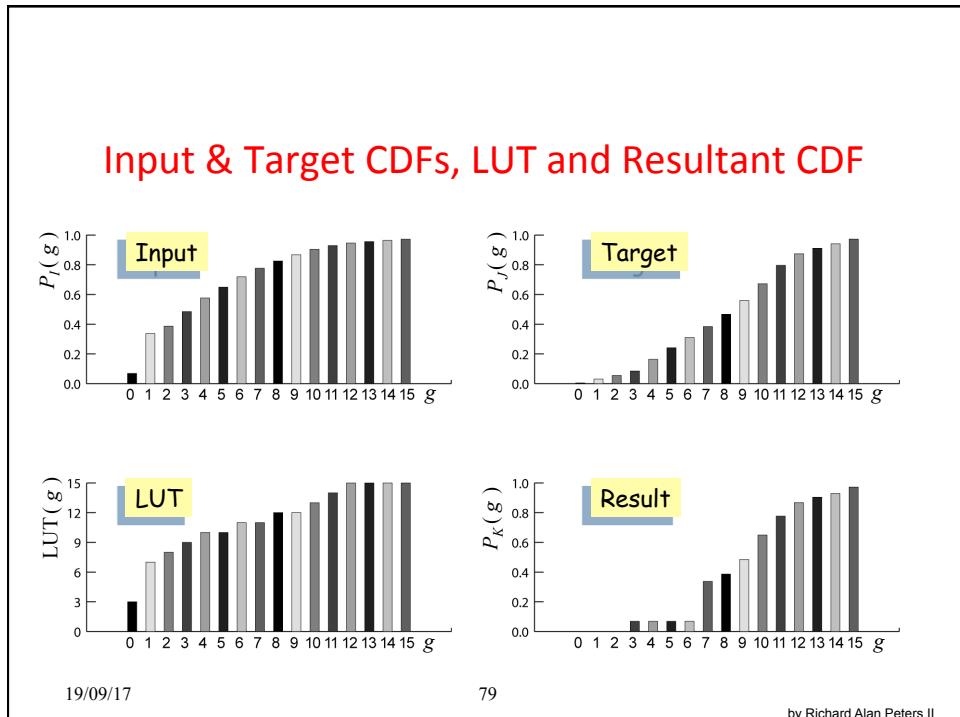
```
K = uint8(LUT(double(I)+1));
```

19/09/17

76

by Richard Alan Peters II





Remap a color Image:

To Have Two of its Color
pdfs Match the Third



original

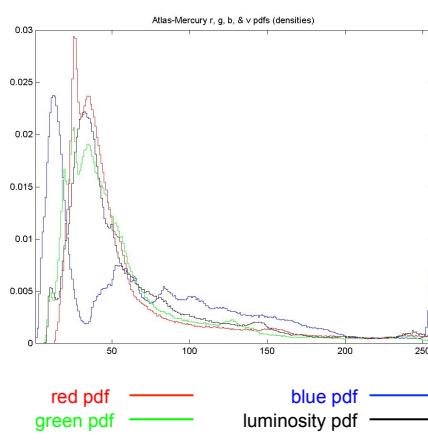
G & B \leftarrow RB & R \leftarrow GR & G \leftarrow B

19/09/17

81

by Richard Alan Peters II

Probability Density Functions of a Color Image

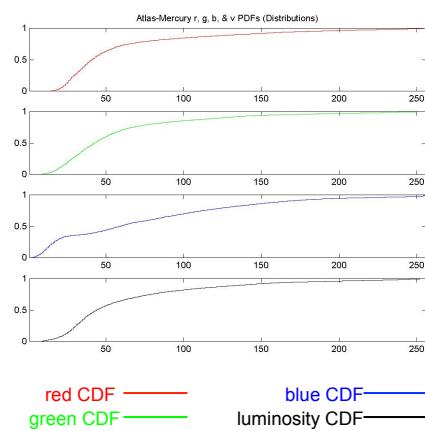


19/09/17

82

1999-2007 by Richard Alan
Peters II

Cumulative Distribution Functions (CDF)

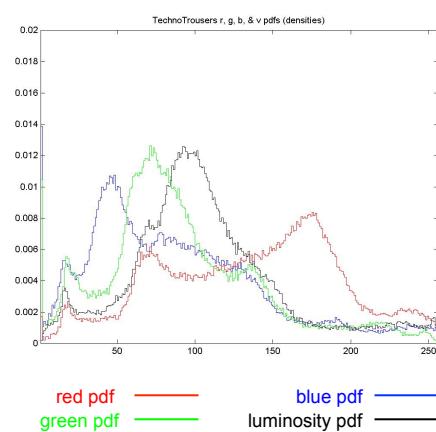


19/09/17

83

1999-2007 by Richard Alan
Peters II

Probability Density Functions of a Color Image

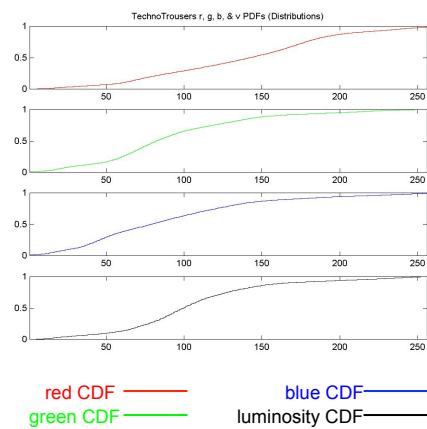


19/09/17

84

1999-2007 by Richard Alan
Peters II

Cumulative Distribution Functions (CDF)



19/09/17

85

1999-2007 by Richard Alan
Peters II

Remap an Image to have the Lum. CDF of Another



original



target



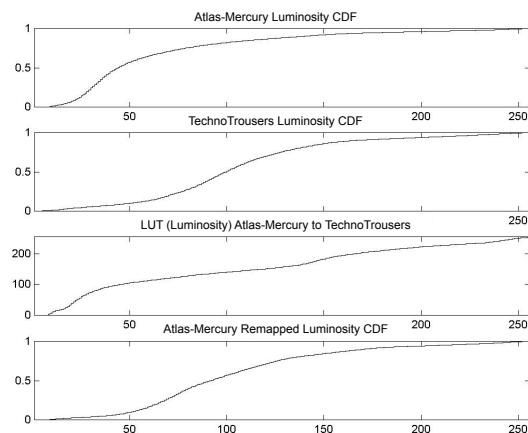
luminosity remapped

19/09/17

86

1999-2007 by Richard Alan
Peters II

CDFs and the LUT

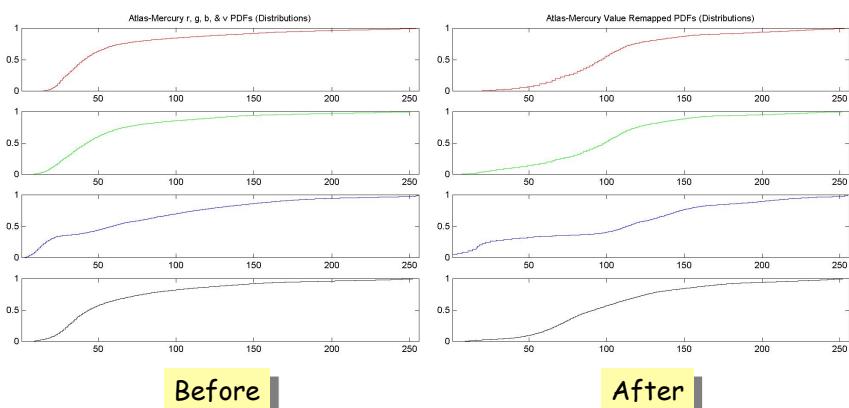


19/09/17

87

1999-2007 by Richard Alan
Peters II

Effects of Luminance Remapping on CDFs



19/09/17

89

1999-2007 by Richard Alan
Peters II

Remap an Image to have the rgb CDF of Another



original



target



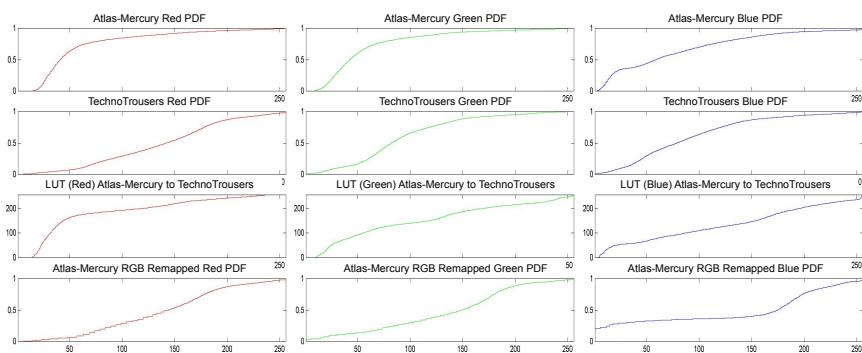
R, G, & B remapped

19/09/17

90

1999-2007 by Richard Alan
Peters II

CDFs and the LUTs



19/09/17

91

1999-2007 by Richard Alan
Peters II

END OF LECTURE

Recall Learning objectives of Week 2: Students will be able to:

1. Discuss the main problems of computer (artificial) vision, its uses and applications
2. Design and implement various image transforms: point-wise transforms

Reading Assignments:

[Klette Book] 1.1, 1.3, 2.1.1, 2.1.2

[Gonzalez and Woods]: Chap 3.1 – 3.3

HW Assignment: you get your hands dirty by starting with your first Image Processing implementations